

MiniProject 3: Classification of Image Data

COMP 551, Winter 2021, McGill University
Contact TAs: Safa Alver and Tianyu Shi

Please read this entire document before beginning the assignment.

Preamble

- This mini-project is **due on March 25th at 11:59pm (EST, Montreal Time)**. There is a penalty of 2^k percent penalty for k days of delay, which means your grade will be scaled to be out of $100 - 2^k$. No submission will be accepted after 6 days of delay.
- This mini-project is to be completed in groups of three. All members of a group will receive the same grade except when a group member is not responding or contributing to the project. If this is the case and there are major conflicts, please reach out to the group TA for help and flag this in the submitted report. Please note that it is not expected that all team members will contribute equally. However every team member should make integral contributions to the project, be aware of the content of the submission and learn the full solution submitted.
- You will submit your assignment on MyCourses as a group. You must register your group on MyCourses and any group member can submit. See MyCourses for details.
- We recommend to use **Overleaf** for writing your report and **Google colab** for coding and running the experiments. The latter also gives access to the required computational resources. Both platforms enable remote collaborations.
- There are additional cloud compute resources available for this project. Please check MyCourses for instructions on how to access those.
- You should use Python for this mini-project. You are free to use libraries with general utilities, such as matplotlib, numpy and scipy for Python, unless stated otherwise in the description of the task. In particular, in most cases you should implement the models and evaluation functions yourself, which means you should not use pre-existing implementations of the algorithms or functions as found in SciKit learn, and other packages. The description will specify this in a per case basis.

Background

In this miniproject, you will **implement a multilayer perceptron from scratch**, and use it to **classify image data**. The goal is to implement a basic neural network and its training algorithm from scratch and get hands-on experience with important decisions that you have to make while training these models.

Task 1: Acquire the data

Your first task is to acquire the image dataset. You will be using only one dataset in your experiments: **MNIST**. Use the MNIST dataset with the default train and test partitions. You can use `tf.keras.datasets.mnist.load_data()` to load the data, see the tutorial <https://www.tensorflow.org/quantum/tutorials/mnist> for more




Figure 1: An MLP with 2 hidden layers each having 4 units.

information. Note that since you will be working with multilayer perceptrons, after loading the data, you will have to vectorize it so that it can have the appropriate dimensions. Also do not forget to normalize the training and test set (see <https://cs231n.github.io/neural-networks-2/#datapre>).

Based on your previous miniprojects, you might be asking the question: where are the features? Well, this is the whole point of using neural nets: instead of hand-designing the features, you train the model so that the feature extractor is also learned together with the classifier on top.

Task 2: Implement a Multilayer Perceptron

In this mini-project, you will implement a multilayer perceptron (MLP) to classify image data. An MLP is composed of three types of layers: (1) an input layer, (2) hidden layers, (3) an output layer (see Figure 1). You should implement it from scratch based on the code available in the slides. Your implementation should include the backpropagation and the mini-batch gradient descent algorithm used (e.g., SGD).

You are free to implement the MLP as you see fit, but you should follow the equations that are presented in the lecture slides, and you must implement it from scratch (i.e., you **cannot** use TensorFlow or PyTorch or any other library). Using the Numpy package is encouraged. Regarding the implementation, we recommend the following approach:

- Implement the MLP as a python class. The constructor for the class should take as input the activation function (e.g., ReLU), the number of hidden layers (e.g., 2) and the number of units in the hidden layers (e.g., [64, 64]) and it should initialize the weights and biases (with an initializer of your choice) as well as other important properties of the MLP.
- The class should have (at least) two functions:
 - A `fit` function, which takes the training data (i.e., \mathbf{X} and \mathbf{y})—as well as other hyperparameters (e.g., the learning rate and number of gradient descent iterations)—as input. This function should train your model by modifying the model parameters.
 - A `predict` function, which takes a set of input points (i.e., \mathbf{X}) as input and outputs predictions (i.e., $\hat{\mathbf{y}}$) for these points.
- In addition to the model classes, you should also define a function `evaluate_acc` to evaluate the model accuracy. This function should take the true labels (i.e., \mathbf{y}), and target labels (i.e., $\hat{\mathbf{y}}$) as input, and it should output the accuracy score.

You are also free to use any Python libraries you like to tune the hyper-parameters, for example see https://scikit-learn.org/stable/modules/grid_search.html.

Task 3: Run the experiments and report

The goal of the experiments in this part is to have you explore the consequences of important decisions made while training neural networks. *Split the dataset into training and test sets. Use test set to estimate performance in all of the experiments after training the model with training set. Evaluate the performance using accuracy.* You are welcome to perform any experiments and analyses you see fit (e.g., the effect of data augmentation / dropout / number of hidden layers / ... on accuracy), **but at a minimum you must complete the following experiments in the order stated below:**

1. First of all, create three different models: (1) an MLP with no hidden layers, i.e., it directly maps the inputs to outputs, (2) an MLP with a single hidden layer having 128 units and ReLU activations, (3) an MLP with 2 hidden layers each having 128 units with ReLU activations. It should be noted that since we want to perform classification, all of these models should have a softmax layer at the end. After training, compare the test accuracy of these three models on the MNIST dataset. Comment on how non-linearity and network depth effects the accuracy. Are the results that you obtain expected?
2. Take the last model above, the one with 2 hidden layers, and create two different copies of it in which the activations are now sigmoid and tanh. After training these two models compare their test accuracies with model having ReLU activations. **Comment on the performances of these models: which one is better and why? Are certain activations better than others? If the results are not as you expected, what could be the reason?**
3. Create an MLP with 2 hidden layers each having 128 units with ReLU activations as above. However, this time, **add L2 regularization (weight decay) to the cost and train the MLP in this way.** How does this affect the accuracy?
4. Create an MLP with 2 hidden layers each having 128 units with ReLU activations as above. **However, this time, train it with unnormalized images.** How does this affect the accuracy?
5. You can report your findings either in the form of a table or a plot in the write-up. However, **include in your colab notebooks the plots of the test and train performance of the MLPs as a function of training epochs.** This will allow you to see how much the network should be trained before it starts to overfit to the training data.

Note: The above experiments are the minimum requirements that you must complete; however, this project is open-ended. For example, you might investigate the width (number of units in the hidden layers) of the MLP on its test accuracy. We would also love to see possible the effect of different regularizations on the final performance. Another interesting thing to report might be training the MLP with 10^k , $k \in \{0, 1, 2, 3, 4\}$ images and plotting the test accuracy. You do not need to do all of these things, but you should demonstrate creativity, rigour, and an understanding of the course material in how you run your chosen experiments and how you report on them in your write-up.

Deliverables

You must submit two separate files to myCourses (**using the exact filenames and file types outlined below**):

1. **code.zip:** Your model implementation, and its training and evaluation code (as some combination of .py and .ipynb files).
2. **writeup.pdf:** Your (max 5-page) project write-up as a pdf (details below).

Write-up instructions

Your team must submit a project write-up that is a maximum of five pages (single-spaced, 11pt font or larger; minimum 0.5 inch margins, an extra page for references/bibliographical content can be used). We highly recommend that students use \LaTeX to complete their write-ups. You have some flexibility in how you report your results, but you must adhere to the following structure and minimum requirements:

Abstract (100-250 words) Summarize the project task and your most important findings.

Introduction (5+ sentences) Summarize the project task, the datasets, and your most important findings. This should be similar to the abstract but more detailed. You should include background information and a few citations to relevant work (e.g., other papers analyzing these datasets).

Datasets (5+ sentences) Very briefly describe the dataset. Present the exploratory analysis you have done to understand the data, e.g. class distribution.

Results (7+ sentences, possibly with figures or tables) Describe the results of all the experiments mentioned in **Task 3** (at a minimum) as well as any other interesting results you find (Note: demonstrating figures or tables would be an ideal way to report these results).

Discussion and Conclusion (5+ sentences) Summarize the key takeaways from the project and possibly directions for future investigation.

Statement of Contributions (1-3 sentences) State the breakdown of the workload across the team members.

Evaluation

The mini-project is out of 100 points, and the evaluation breakdown is as follows:

- Completeness (20 points)
 - Did you submit all the materials?
 - Did you run all the required experiments?
 - Did you follow the guidelines for the project write-up?
- Correctness (40 points)
 - Are your models implemented correctly?
 - Are your reported accuracies close to our solution?
 - Do you observe the correct trends in the experiments (e.g., how the accuracy changes as the depth of the MLP increases)?
 - Do you observe the correct impact of activation choice, regularization and normalization on the model performance?
- Writing quality (30 points)
 - Is your report clear and free of grammatical errors and typos?
 - Did you go beyond the bare minimum requirements for the write-up (e.g., by including a discussion of related work in the introduction)?
 - Do you effectively present numerical results (e.g., via tables or figures)?
- Originality / creativity (10 points)
 - Did you go beyond the bare minimum requirements for the experiments?
 - **Note:** Simply adding in a random new experiment will not guarantee a high grade on this section! You should be **thoughtful and organized** in your report.

Final remarks

You are expected to display initiative, creativity, scientific rigour, critical thinking, and good communication skills. You don't need to restrict yourself to the requirements listed above - feel free to go beyond, and explore further.

You can discuss methods and technical issues with members of other teams, but **you cannot share any code or data with other teams.**